

Markov Model, HMM

Special case of Bayes Net? Yes and No. Joint distribution $P(X_0, \dots, X_T) = P(X_0) \prod_t P(X_t | X_{t-1})$

k-th model allow dependencies on k earlier steps.

Forward Algorithm: $P(X_t) = \sum_{x_{t-1}} P(X_t | X_{t-1}=x_{t-1})$
 $= \sum_{x_{t-1}} P(X_{t-1}=x_{t-1}) P(X_t | X_{t-1}=x_{t-1})$

$$P(X_0, X_1, \dots, X_T, E_T) = P(X_0) \prod_{t=1:T} P(X_t | X_{t-1}) P(E_t | X_t)$$

$$B|X_t| = P(X_t | e_{1:t}) \Rightarrow P(X_{t+1} | e_{1:t}) = \sum_{x_t} P(X_{t+1} | x_t) B(x_t) \text{ Elapse time}$$

$$B'|X_{t+1}| = P(X_{t+1} | e_{1:t}) \Rightarrow P(X_{t+1} | e_{1:t}) \propto_{x_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) = P(e_{t+1} | X_{t+1}) B'(X_{t+1}) \text{ observe}$$

Cost per time step: $O(|X|^2)$, X is the number of states **Filtering (Forward)**

Viterbi Algorithm (max)

For each state at time t, keep track of the (unweighted) **maximum probability of any path** to it:

$$m_{1:t}(x_t) = \max_{x_{t-1}} P(x_t | x_{t-1})$$

$$m_{1:t+1} = \text{VITERBI}(m_{1:t}, e_{t+1}) \\ = P(e_{t+1} | X_{t+1}) \max_{x_t} P(X_{t+1} | x_t) m_{1:t}(x_t)$$

Forward Algorithm (sum)

For each state at time t, keep track of the **total probability of all paths** to it:

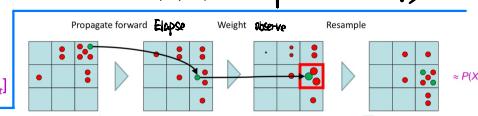
$$f_{1:t}(x_t) = P(x_t | e_{1:t}) \\ = \sum_{x_{t-1}} P(x_t | x_{t-1}) f_{1:t-1}(x_{t-1})$$

$$f_{1:t+1} = \text{FORWARD}(f_{1:t}, e_{t+1})$$

$$= \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(X_{t+1} | x_t) f_{1:t}(x_t)$$

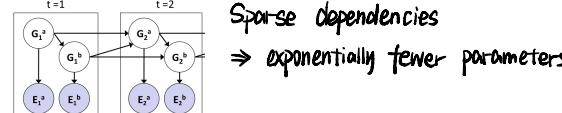
Most likely explanation = most probable path

Time: $O(|X|^2 T)$ Space: $O(|X|T)$



DBN: repeat a fixed Bayes Net structure at each time

Inference: variable elimination. "unroll" the network



|X| is too big to exact inference. particle filtering
 $w|x\rangle = P(e|x)$. $B(X) \propto P(e|X) \cdot B'(X)$

MDPs optimal policy: $\pi^*: S \rightarrow A$. given action from state. maximize except utility.

An MDP is defined by:

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
- A start state
- Maybe a terminal state

Value Iteration $O(S^2 A)$

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Policy extraction

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Policy Iteration

Step 1: Policy evaluation: calculate utilities for some fixed (not optimal) policy

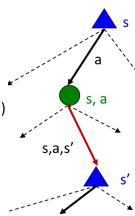
Step 2: Policy improvement: update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values

In value iteration:

- Every iteration updates both the values and (implicitly) the policy
- We don't track the policy, but taking the max over actions implicitly recomputes it

In policy iteration:

- We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
- After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
- May converge faster



s is a state
 $V^*(s) = \text{expected utility starting in } s \text{ and acting optimally}$

(s, a) is a q-state
 $Q^*(s, a) = \text{expected utility starting out having taken action } a \text{ from state } s \text{ and (thereafter) acting optimally}$

(s, a, s') is a transition
 $T(s, a, s') = P(s' | s, a), R(s, a, s')$

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Idea 1: Iterative updates (like value iteration) **Evaluation**

- Start with $V_k^*(s) = 0$
- Given V_k^* , calculate the depth k+1 values for all states:

$$V_{k+1}^*(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^*(s')]$$
- Repeat until convergence
- Efficiency: $O(S^2)$ per iteration

Idea 2: Without the maxes, the Bellman equations are just a linear system

$$V^*(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^*(s')]$$

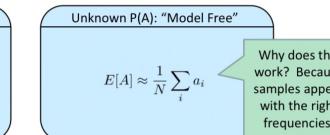
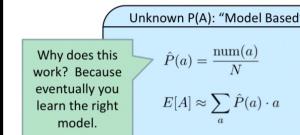
$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \text{ Improvement}$$

All use one-step lookahead expectimax fragments.

Reinforcement Learning a MDP, don't know T or R

Model - Based: learn empirical MDP model, solve it.

Without P(A), instead collect samples $[a_1, a_2, \dots, a_n]$



Model - Free Passive

Direct Evaluation: Average together observed sampled value.

Active Q - learning

$$\text{sample} = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

$$Q(s, a) \leftarrow (1 - \alpha) Q(s, a) + (\alpha) [\text{sample}]$$

Explore Idea: select actions based on modified Q-value

- Exploration function: takes a Q-value estimate u and a visit count n, and returns an optimistic utility, e.g.
 $f(u, n) = u + k/n$

Feature - Based Representation

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

Policy Search: start with an OK solution.

fine-tune feature weights to find better policy

Sample of V(s):

$$\text{sample} = R(s, \pi(s), s') + \gamma V^\pi(s')$$

Update to V(s):

$$V^\pi(s) \leftarrow (1 - \alpha) V^\pi(s) + (\alpha) \text{sample}$$

Same update:

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha (\text{sample} - V^\pi(s))$$

Exponential moving average

The running interpolation update: $\bar{x}_n = (1 - \alpha) \cdot \bar{x}_{n-1} + \alpha \cdot x_n$

$$\bar{x}_n = \frac{x_n + (1 - \alpha) \cdot x_{n-1} + (1 - \alpha)^2 \cdot x_{n-2} + \dots}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots}$$

Makes recent samples more important

Forgets about the past (distant past values were wrong anyway)

transition = (s, a, r, s')

$$\text{difference} = [r + \gamma \max_{a'} Q(s', a')] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

Exact Q's

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$

Approximate Q's

Supervised Machine Learning

Overfit Sol: more training data (less sampling variance, training more like test)

limit the complexity of hypotheses

(regularization or small hypothesis space)

To generalize better: smooth or regularize the estimate

empirical rate $P_{ML}(x) = \frac{\text{count}(x)}{\text{total samples}}$ \Leftrightarrow Maximum Likelihood

most likely parameter $\theta_{MAP} = \arg \max_\theta P(\theta | x)$

Laplace Smoothing $= \arg \max_\theta P(x | \theta) P(\theta)$

learn an unknown target function f

$$P(Y, W_1 \dots W_n) = P(Y) \prod_i P(W_i | Y)$$

Word at position i, not i^{th} word in the dictionary!

Usually, each variable gets its own conditional probability distribution $P(F | Y)$

Here

- Each position is identically distributed

All positions share the same conditional probabilities $P(W | Y)$

This is called "bag-of-words" because model is insensitive to word order or reordering

$$\theta_{ML} = \arg \max_\theta P(X | \theta)$$

$$= \arg \max_\theta \prod_i P(x_i | \theta)$$

Another option: linear interpolation

Also get the empirical $P(X)$ from the data

Make sure the estimate of $P(X | Y)$ isn't too different from the empirical $P(X)$

$$P_{LIN}(x | y) = \alpha \hat{P}(x | y) + (1.0 - \alpha) \bar{P}(x)$$

$P_{LAP}(x) = \frac{c(x) + 1}{\sum_x c(x) + 1} \xrightarrow{k \text{ times}} P_{LAP,k}(x) = \frac{c(x) + k}{N + k |X|}$
 $c(x) + 1$
 $N + |X|$
 k is strength of prior. $k \uparrow$, fit \downarrow

$P_{LAP,k}(x | y) = \frac{c(x, y) + k}{c(y) + k |X|}$
Condition

Linear Classifiers activation_w(x) = $\sum_i w_i \cdot f_i(x) = w \cdot f(x)$ (+ bias = constant feature) adjust: $w' = w + y^* \cdot f$

Probabilistic Decision $\rightarrow \phi(z) = \frac{1}{1+e^{-z}}$

$$\max_w ll(w) = \max_w \sum_i \log P(y^{(i)}|x^{(i)}; w)$$

$$P(y^{(i)} = +1|x^{(i)}; w) = \frac{1}{1 + e^{-w \cdot f(x^{(i)})}}$$

$$P(y^{(i)} = -1|x^{(i)}; w) = 1 - \frac{1}{1 + e^{-w \cdot f(x^{(i)})}}$$

logistic regression

$$z_1, z_2, z_3 \rightarrow \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3}}, \frac{e^{z_2}}{e^{z_1} + e^{z_2} + e^{z_3}}, \frac{e^{z_3}}{e^{z_1} + e^{z_2} + e^{z_3}}$$

original activations

$$\max_w ll(w) = \max_w \sum_i \log P(y^{(i)}|x^{(i)}; w)$$

multiclass

$$P(y^{(i)}|x^{(i)}; w) = \frac{e^{w_y \cdot f(x^{(i)})}}{\sum_y e^{w_y \cdot f(x^{(i)})}}$$

consider: $g(w_1, w_2)$

Updates: $w_1 \leftarrow w_1 + \alpha * \frac{\partial g}{\partial w_1}(w_1, w_2)$

$$w_2 \leftarrow w_2 + \alpha * \frac{\partial g}{\partial w_2}(w_1, w_2)$$

Updates in vector notation:

$$w \leftarrow w + \alpha * \nabla_w g(w)$$

with: $\nabla_w g(w) = \begin{bmatrix} \frac{\partial g}{\partial w_1}(w) \\ \frac{\partial g}{\partial w_2}(w) \end{bmatrix}$ = gradient

Gradient Ascent

Batch

$$w \leftarrow w + \alpha * \sum_i \nabla \log P(y^{(i)}|x^{(i)}; w)$$

$$P(y^{(i)}|x^{(i)}; w) = \frac{e^{w_y \cdot f(x^{(i)})}}{\sum_y e^{w_y \cdot f(x^{(i)})}}$$

add f to weights of the correct class

Increase the score of the correct class

Stochastic

Idea: once gradient on one training example has been computed, might as well update before computing next one

```

    init w
    for iter = 1, 2, ...
        pick random j
        w ← w + α * ∇ log P(y(j)|x(j); w)
    
```

Mini - Batch

Idea: gradient over a small set of training examples (=mini-batch) can be computed in parallel, might as well do that instead of a single one

```

    init w
    for iter = 1, 2, ...
        pick random subset of training examples J
        w ← w + α * ∇ log P(y(j)|x(j); w)
    
```

Batch GD
 -Slowest
 -Perfect gradient
Stochastic GD
 -Fastest
 -Rough-estimate grad
Mini-batch GD
 -Compromise

L2 loss function: sum of squared errors over all examples

$$L(\mathbf{w}) = \sum_i (y_i - h_{\mathbf{w}}(\mathbf{x}_i))^2 = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

We want the weights \mathbf{w}^* that minimize loss

Analytical solution: at \mathbf{w}^* the derivative of loss w.r.t. each weight is zero

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

\mathbf{X} is the data matrix (all the data, one example per row); \mathbf{y} is the vector of output values

Multiclass Perceptron

Start with all weights = 0

Pick up training examples one by one
 Predict with current weights

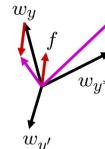
$$y = \arg \max_y w_y \cdot f(x)$$

If correct, no change!

If wrong: lower score of wrong answer, raise score of right answer

$$w_y = w_y - f(x)$$

$$w_{y^*} = w_{y^*} + f(x)$$



Unsupervised

k-means non-deterministic

$$\phi(\{x_i\}, \{a_i\}, \{c_k\}) = \sum_i \text{dist}(x_i, c_{a_i})$$

points

assignments

squared Euclidean

distance

update assignments of means

Cannot increase ϕ

$$-|\mathbf{x} - \mu|^2$$

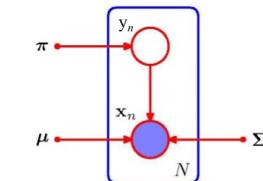
$$p(\mathbf{x}|\mu, \Sigma) = \frac{1}{\sqrt{2\pi}^d} e^{-\frac{|\mathbf{x} - \mu|^2}{2\Sigma^2}}$$

P(Y): Distribution over k components (clusters)

P(X|Y): Each component generates data from a **Multivariate Gaussian** with mean μ_i and covariance matrix Σ_i

Each data point is sampled from a **generative process**:

1. Choose component i with probability π_i
2. Generate data point from $N(\mathbf{x}|\mu_i, \Sigma_i)$



[E step] Compute probability of each instance having each possible label

[M step] Treating each instance as fractionally having both labels, compute the new parameter values

EM degrades to k-means if we assume

- All the Gaussians are spherical and have identical weights and covariances
 - i.e., the only parameters are the means
- The label distributions computed at E-step are point-estimations
 - i.e., hard-assignments of data points to Gaussians
 - Alternatively, assume the variances are close to zero

Overfitting is also possible in regression

Regularization can be used to alleviate overfitting

LASSO (Least Absolute Shrinkage and Selection Operator)

$$L(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_k |w_k|$$

Ridge Regression

$$L(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_k w_k^2$$